

COM-480 Data Visualization

Process book

Auguste Baum, Yanis Berkani & Clément Petit

June 2, 2021



1 Dataset & Exploratory Data Analysis

In such a project, choosing a clean dataset and conducting a thorough Exploratory Data Analysis (EDA) is very important. Indeed, in order to have good visualizations, we need to have useful and consistent data, as well as the easiest access to information we can get. To achieve this result, a first step is to spend some time looking for a pretty clean and interesting dataset about the desired subject, which in our case is music. After searching for some time, we got our data from Spotify, cleaned for a Kaggle competition.¹ Then, before getting to the heart of the matter, comes the part of understanding and analyzing the dataset(s). Finally, after that we set some precise goals towards the main objective: representing information in a visually interesting way.

For instance, in our case we decided to merge two datasets (tracks and artists), to have access to both information about the track and the artist in a single dataset. Then, we handled strange track names, we removed completely silent tracks, we dropped duplicates and so on. Finally, we processed the data to best fit our project's main goals. For example, one of them is to visualize the evolution of popular genres, hence we computed a dataframe containing the top 10% songs in terms of popularity along with their genre for each year.

All of this takes a lot of time as we never know how precise we should get, and when we have something in mind, even small fixes, it's very hard to let go. One of the choices that we made to make our task way easier when dealing with multi-artist

¹<https://www.kaggle.com/yamaerenay/spotify-dataset-19212020-160k-tracks>

tracks was to only take the first one appearing in the featuring. We tried other solutions but it this seemed to have the best balance between complexity of the method and usability of the result.

Note that you can find out a lot more details about the features and look at our data cleaning and processing in our repository,² in the notebook called `EDA.ipynb`.

Let us now describe our process for the visualisations.

2 Visualisations – Design and Implementation

2.1 Timeline visualisation

The timeline visualisation gives an overview of the most popular tracks through the years, since the 1920s. We plot the most popular song for each year in our dataset. Inspired by the lectures and exercises, we had the idea to incorporate this visualisation in a “Brush & Zoom” chart. After selecting one song for each year between 1922 and 2021, there were a few difficulties we had to think about:

- how to visually represent the songs;
- how to dynamically add or remove songs to the graph, depending on the zoom amount;
- how and when to display the song information;
- how to make the zoom and brush behave well in this particular case as we cannot zoom indefinitely. The zoom has to stop when each consecutive year is visible on the x -axis. Hence it should not be allowed to rescale the graph directly with the Brush component.

After some research and several attempts, we came up with a first proof of concept, shown in Figure 1. The songs are represented by green circles, with year on the x -axis and popularity on the y -axis. Hovering on a circle causes it to be replaced by a tooltip containing the song title and artist.




Figure 1: A proof-of-concept of the timeline visualisation. The figure on the right shows the tooltip displayed on hover, containing track information.

²<https://github.com/com-480-data-visualization/data-visualization-project-2021-vizbrains>

Then, we had the idea to incorporate sound samples for the user to have a glimpse at each song. After all, in a project about music, it makes sense to incorporate some sound elements. Spotify offers this possibility with their “Spotify Embed” widget (see [Figure 2](#)). At first the idea was to include the sample directly in the tooltip when hovering on a song, but we quickly realised that our implementation with D3.js made this impractical.




Figure 2: Spotify Embed widget, in compact format.

Another approach was to display the Spotify widget on a fixed position, at the top of the visualisation. We finally went with this solution. When the cursor is moved from the song, the sample stays at the top of the visualisation, until another song is hovered over. Furthermore, the songs titles and artists being displayed automatically in the Spotify widget, there was no longer any point in displaying this information in tooltips so we removed these. Instead, we added a small hovering effect on the song circles. The final result is showed on [Figure 3](#). As a final note on this first visualisation, regarding the colors, we chose to stay simple, and pretty much match the Spotify color theme.

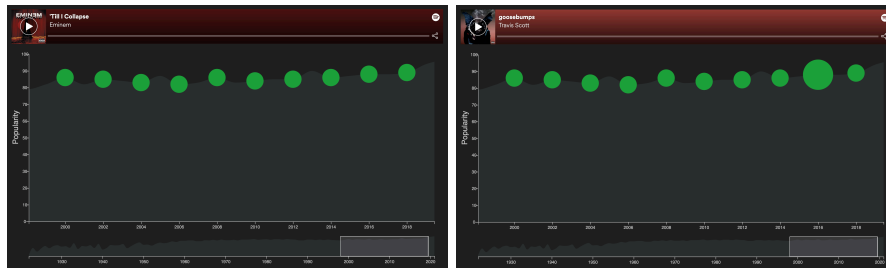


Figure 3: Final version of the Timeline visualization.

2.2 Year comparator visualisation

The bubble visualisation aims to compare 2 years, according to some of the features present in the dataset. It is still a temporal exploration of the data which is similar to the timeline, except for that it is not an overview anymore (but just a comparator). The intended use-case was for example to give insight about claims like: “music is generally easier to dance to now than during the disco years” or “music generally gave more good vibes during the Beatles era”.

The data plotted is the yearly average of the 20% most popular songs in the dataset, where the values are normalised to the $[0, 1]$ interval.

Coming from a scientific background, the most obvious way for us to achieve this would have been a bar chart with the features on the x axis and the values on the y axis. Though this might be an efficient way to visualise this, we were sure that the freedom that comes with the Web would lead us to a more innovative and powerful way to display our findings.

As a first step, we considered using *bubbles* of varying area. The choice of bubbles is often made on map visualisations to show where effects are most prominent, but we argue this is often a poor choice, using the visualisation on [this Johns Hopkins university website](#) as an example:

- the amount of bubbles often makes the visual very messy;
- the amount of bubbles can make the visual very computationally intensive: the above web page takes **20 to 30 seconds** from load time to displaying the bubbles;
- humans are often not good at perceiving areas (see lecture 7.2.), which makes comparison difficult.

For maps, we maintain that coloring is a more sensible choice.

However, in our case we avoided these pitfalls: we superimposed the bubbles so that we could clearly see which was the biggest, and the absolute value of the features was secondary to us so we did not have to worry about issues with perception. You can see our first prototype in [Figure 4](#).



Figure 4: The first version of the bubble year comparator visualisation. The red bubbles show the average of each feature for 1950 and the green for 2001.

One can see that for each feature, the smallest bubble is always in the foreground, which is obviously desirable otherwise we might only see one bubble. This was somewhat technically challenging, because the concept of z -order (i.e. foreground and background) does not really exist in SVG: it's the order in which the elements are placed that determines this. We managed to solve the issue with use tags, which make a copy of an element.

At a glance this visualisation might seem satisfactory, but one can notice an important issue with this. What we want to show is when the years differ, but in [Figure 4](#), both danceability and valence are about the same between the two years, and yet the two bubbles are completely different colors, which is incoherent.

An alternative was to have the front bubble be 50% transparent, so that when the two bubbles overlap the front color is always the same, as can be seen in [Figure 5](#). But in this case we now have 3 colors, the front color that stays the same regardless of which bubble is smallest, the back green color and the back red color (that can be seen on the outside of the Valence bubble). This is also disorienting for the user.




Figure 5: The same as [Figure 4](#) but the front bubble is 50% transparent.

Finally, updating the size of the bubbles dynamically proved to be a bit too difficult, especially regarding those use tags that were mentioned earlier, and that were used to put the smallest bubble in front.

For all these reasons, we decided to change the visualisation format.

Our second (and current) version is a “lollipop” graph that resembles a scatter plot. [Figure 6](#) shows the visualisation in its current state, showing the exact same data that was used in [Figure 4](#) and [Figure 5](#).

With this format comparison is also very easy, but now the absolute value of the bubbles can also be displayed and understood easily.

Besides, the issue of having a certain element in the foreground is no longer relevant.

The user of course has the choice of both years they want the data for. The way this was done was with sliders. The rationale behind this and the technical implementation are detailed in [subsection 3.1](#).

One debate was whether or not the user should also be given the choice of the features that are displayed. Indeed, who is to say that they want to know about acoustiness and not about danceability? Adding as much interaction as possible might make sense as a way to reinforce engagement, but we are also aware (from experience) that too much involvement from the user can lead to fatigue and loss of interest. Besides, some features tend to vary only a little compared to others that cover the whole spectrum, so these might give




Figure 6: The bubble visualisation in its current state.

the impression that our visualisation is uninteresting. For these reasons, we opted to choose the best features and make this choice final. Of course if the user knows Javascript they could go change the features in the source code!

2.3 Genres popularity visualisation

The last visualization consists in depicting music genre trends over the last century. To be more precise, it is a dynamic bar chart showing the popularity of the main music genres for a chosen year.

We directly thought of the idea of studying the evolution of musical tastes, but we did not know how to visually represent it at first. After some research, we were reminded of the spectrum technique that is often used to visualise musical features in *individual* tracks (see [Figure 7](#)), which lead us to the interactive bar chart idea.



Figure 7: Music spectrum design that inspired us for the bar graph.

However there are some aspects that we had not anticipated at first:

- **Non uniform distribution:** Genres popularity is unfortunately not distributed as uniformly as in [Figure 7](#), hence creating such a harmonious musical wave is impossible in our case.
- **Genres ordering:** At first we wanted the graph to always keep the bars in descending order of popularity, for the user to instantly see what are the main genres of a particular year. However, in the end, we decided to rather order it randomly as it looks more like the spectrum and is still clear. The user still has the option to order the graph for the current year by clicking on a button. It is even more interesting as this way, they can also follow the evolution of top genres of a certain year throughout the century.
- **Transition:** The idea that we mentioned in Milestone 2 (see [Figure 8](#)) was to have a “loading” phase that would look similar to [Figure 7](#). However, we finally decided to simply transition from previous to current state, as it actually looked nicer.

3 Individual features

3.1 Sliders

The sliders can be used to change a given parameter dynamically so that the change can happen as the slider is moved. This is not quite the same as a drop-down menu

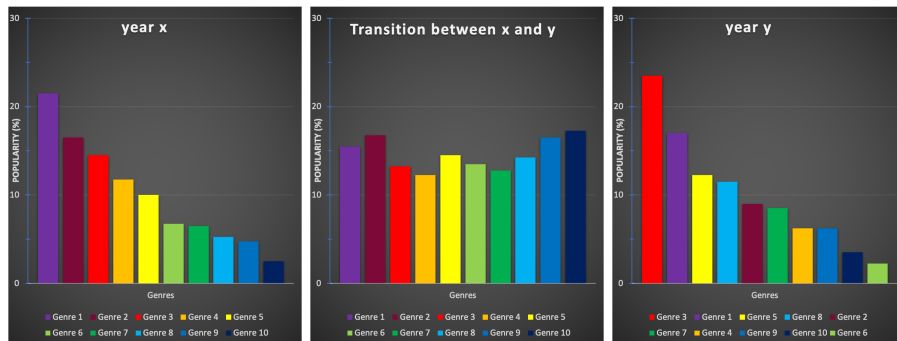


Figure 8: A schematic view of the transition between years for the bar chart as we imagined it in milestone 2.

because the former can be used to get a general idea of the behaviour of the data as the parameter changes, without having to plot the data directly.

For example, the main purpose of the bubble visualisation is to compare two years; the timeline visualisation is obviously best for visualising the behaviour as the year changes, but plotting 5 curves at once would make it unreadable. The slider feature in the bubble visualisation is a compromise between the number of points of comparison and the temporal behaviour.

This feature involved refreshing the data that is shown, and hence it was made easier using D3's `enter` and `update` constructs. However, one issue is that it seems difficult to avoid repeated code. Indeed, the drawing has to be made once at load time for the `circle` elements to be created, (see the `draw` method) and then it has to be refreshed using the new data at each new slider position; a `transition()` call is used then to guarantee fluidity and clarity. These two phases, though extremely similar, did not seem to be modularisable.

3.2 Tooltips

The idea of tooltips is to only show information when it is needed; it allows for different levels of precision in a visualisation, so that it can reach a wider range of people. Indeed, some might be happy with a less detailed but more elegant picture, whereas others might prefer sacrificing a bit of simplicity to get more precise information.

We use tooltips in several places throughout the project:

- in the timeline visualisation we can display the most popular song of each year;
- in the bubble visualisation we give a description of the musical features being plotted—though most are self-explanatory, some can appear a bit cryptic.

First we used `rect` elements that were placed and removed on mouse events, but then we used a `div` tag that is made opaque or transparent.

This second attempt, though more efficient (as it doesn't require deleting shapes), proved more difficult than expected because it involves including HTML elements

into an SVG canvas. For example, the placement of the tooltip in the bubble visualisation should be next to the feature that is hovered on, but the `mouseover` function does not have access to the coordinate system of the SVG canvas; we could make those functions internal to the plot constructor, but then the `this` object in `mouseover` and `mouseout` would not refer to the right element anymore.

One major difficulty we found while working on this project was to modularise our code; even though we would much rather have used the same code in both cases, the two use cases were different enough that it seemed not worth the effort.

4 Peer Assessment

Despite the current situation, the three of us managed to collaborate well together on this project. We relied on a good communication, with bi-weekly Zoom calls (on Mondays and Fridays) and a group chat.

At first, a big part of the work was to brainstorm and share ideas, in order to make decisions and to set our goals.

Our final result mainly consists in three visualizations, that we mostly carried out individually. While working in parallel, we shared implementation ideas and code snippets to help each other and maximize efficiency.

Hence we each mostly focused ourselves on the following tasks:

- Clément Petit: Timeline visualisation & Website Design
- Auguste Baum: Year Comparator visualisation & Screen Cast
- Yanis Berkani: Genres Popularity visualisation & EDA

Auxiliary work consisted in writing the reports, for which we all participated equally.

5 Conclusion

We really enjoyed working together as a team on this project and we are very satisfied about the final result. This work allowed us to discover new insights on a topic we all are passionate about, while learning how to use very important web development and data visualization tools.

Finally, we would like to thank our Professor Laurent Vuillon for this course and his help on Friday exercise sessions.